Technical Debt in Latin American Academic Software: A Matched-Sample Comparison with Open-Source Controls

Isaac David Torres-Paredes*, Juan Pablo Aviles-Esparza†,
and Italo Javier Tenempaguay-Granizo‡
Software Engineering, Faculty of Informatics and Electronics
Escuela Superior Politécnica de Chimborazo (ESPOCH), Riobamba 060155, Ecuador
{*isaac.torres, juan.aviles, italo.tenempaguay*}*@espoch.edu.ec*

Abstract

Software produced in university courses and capstone projects is widely assumed to carry more technical debt than comparable open-source work, but this has rarely been tested for Latin America. We compare 274 academic GitHub repositories (identified by thesis and course keywords, with contributors located in Latin America) against 275 non-academic repositories matched by language and size on nine static metrics, using Mann-Whitney U tests with Holm correction and Cliff's delta. Contrary to our hypothesis, academic code shows *significantly lower* cyclomatic complexity (median 2.16 vs. 3.12, $\delta = -0.37$, $p < 10^{-11}$). Inheritance depth is higher in academic Java code ($\delta = +0.29$, $p = 0.003$); coupling, cohesion, duplication, and test-file presence do not differ significantly. A check using stars as a maturity signal does not support attributing this to the small, unlabeled control repositories. SonarQube on a 415-repository subset corroborates the complexity result ($\delta = -0.458$, $p < 10^{-14}$) and finds a slightly lower technical debt ratio in academic code ($\delta = -0.142$, $p = 0.025$). Self-admitted technical debt markers run in the hypothesized direction: 42.2% of academic repositories contain at least one, against 23.0% of controls ($\delta = +0.190$, $p < 0.001$).

Keywords: Technical debt, software engineering education, empirical software engineering, Latin American software, code complexity, GitHub mining.

Resumen. Se suele suponer que el software producido en cursos universitarios y trabajos de titulación acumula más deuda técnica que proyectos de código abierto comparables, pero esta suposición casi no se ha contrastado para América Latina. Comparamos 274 repositorios académicos de GitHub (identificados por palabras clave de tesis y cursos, con contribuyentes ubicados en América Latina) con 275 repositorios no académicos emparejados por lenguaje y tamaño, en nueve métricas estáticas, mediante pruebas U de Mann-Whitney con corrección de Holm y delta de Cliff. Contra nuestra hipótesis, el código académico presenta una complejidad ciclomática significativamente *menor* (mediana 2,16 frente a 3,12; $\delta = -0{,}37$; $p < 10^{-11}$). La profundidad de herencia es mayor en el código académico en Java ($\delta = +0{,}29$; $p = 0{,}003$); acoplamiento, cohesión, duplicación y presencia de archivos de prueba no difieren significativamente. Una comprobación con las estrellas como señal de madurez no respalda atribuirlo a los repositorios de control, pequeños y sin etiquetar. SonarQube, aplicado a 415 repositorios, corrobora el resultado de complejidad ($\delta = -0{,}458$; $p < 10^{-14}$) y encuentra un ratio de deuda técnica ligeramente menor en el código académico ($\delta = -0{,}142$; $p = 0{,}025$). Los marcadores de deuda técnica autoadmitida siguen la dirección de la hipótesis: el 42,2 % de los repositorios académicos contiene al menos uno, frente al 23,0 % del grupo de control ($\delta = +0{,}190$; $p < 0{,}001$).

Palabras clave: Deuda técnica, educación en ingeniería de software, ingeniería de software empírica, software latinoamericano, complejidad de código, minería de GitHub.

*Corresponding author: isaac.torres@espoch.edu.ec. ORCID: 0009-0001-7057-9316
†ORCID: 0009-0007-0058-8069
‡ORCID: 0009-0001-5753-4279

1 Introduction

It is commonly assumed, among educators and practitioners alike, that code written for a course assignment or a thesis project carries more technical debt, in the sense of the metaphor introduced by Cunningham [1] (higher complexity, near-zero test coverage, more duplicated code), than comparable work produced outside an academic setting. This assumption motivates a large share of software-engineering curriculum design, yet it is rarely tested with a matched, quantitative comparison, and to our knowledge no such comparison has been run specifically for Latin American academic software, where institutional resourcing, course structure, and open-source contribution culture may all differ from the North American and European settings most existing studies draw on.

This paper runs that comparison directly. We ask whether repositories produced in Latin American university courses and capstone/thesis projects differ measurably, on a set of standard static metrics, from a size- and language-matched sample of non-academic repositories, and if so, in which direction and by how much. We report the result even where it contradicts our starting hypothesis, because a null or reversed finding here is itself informative about what "academic code" means empirically in this region, as opposed to what it is assumed to mean.

The contributions of this work are: (1) a size- and language-matched quantitative comparison of Latin American academic software against non-academic controls, at a scale (549 repositories) sufficient for non-parametric hypothesis testing; (2) a report of an inverted finding, lower and not higher complexity in academic code, with the alternative explanations we could test (language, repository age, size, control-group maturity) and their results; (3) a cross-tool consistency check of the central finding with SonarQube Community, which agrees on the complexity result and isolates a tool- and sample-dependent difference in duplication rate.

2 Related Work

An early in-depth exploration of technical debt by Tom, Aurum, and Vidgen [2] combined a review of the literature with practitioner interviews to describe the dimensions of technical debt (among them code, design and architecture, testing, and documentation debt), why it arises, and its consequences; this vocabulary frames the kinds of debt our metrics and SATD markers can and cannot capture. Li, Avgeriou, and Liang [3] subsequently mapped 94 studies on technical debt and its management, classified TD into ten types, and found that code-related debt has received the most attention, while calling for more empirical studies with high-quality evidence. We are not aware of a prior study that compares academic and non-academic code on these metrics, controlled by language and size, at the scale reported here.

Santos et al. [4] conduct a systematic literature review of external quality metrics for object-oriented software; we apply metrics of the coupling, cohesion, and complexity families they review to a labelled academic/non-academic comparison. More directly relevant, Hamer et al. [5] use Git-derived metrics to measure students' and teams' code-contribution patterns within university software-development courses, across eight projects with 150 students, showing that repository-level metrics can characterize academic software engineering practice; our study extends this direction from process metrics (who contributed what) to product metrics (what the resulting code looks like structurally), and from within-course comparisons to a labeled comparison against non-academic repositories. Mateos et al. [6] apply the same Chidamber and Kemerer [7] object-oriented metric suite this paper uses, in their case to relate service-interface complexity to the underlying object-oriented implementation; their use of CK metrics as an early, development-time complexity signal is consistent with how this paper treats complexity metrics generally, as a proxy available before a system reaches production, not as an after-the-fact quality judgment. The repository-mining methodology follows the practices established by Kalliamvakou et al. [8] and refined by Munaiah et al. [9], both of which show that many GitHub repositories are personal, inactive, or not engineered software, so samples need filtering before structural metrics can be read as representative of software engineering practice; our fork exclusion and 10-commit minimum follow that advice in a simplified form. Munaiah et al. explicitly list homework assignments as an example of the non-engineered repositories that such filtering removes, which is a reminder that our academic sample is, by design, drawn from the kind of repository that engineered-project studies exclude.

The self-admitted technical debt analysis in Section 3.4 is grounded in a body of work that began with Potdar and Shihab [10], who identified SATD in the source-code comments of four large open-source projects (Eclipse, Chromium OS, Apache HTTP Server, and ArgoUML), found it in 2.4% to 31% of files, and reported that more experienced developers introduce most of it, while time pressure and code complexity do not correlate with how much SATD is present; their figures are per file, so they are not directly comparable with our repository-level rates (42.2% academic vs. 23.0% control). Bavota and Russo [11] extend this

line of work to a large-scale study of how SATD spreads and evolves across many open-source projects. Maldonado and Shihab [12] read more than 33,000 comments and classify SATD into five types (design, defect, documentation, requirement, and test debt), with design debt the most common; our marker list (`TODO`, `FIXME`, `HACK`, `XXX`, `BUG`) is a simpler keyword approximation of the comments they analyse manually. Huang et al. [13] subsequently detect SATD in source comments automatically with a text-mining approach that combines feature selection with a composite classifier trained on several projects; we use the simpler keyword approach because precision is less critical for our research question (is the *prevalence* of any SATD marker different between academic and control repositories?) than for studies aimed at classifying SATD type. Alfayez et al. [14] analyse 128,897 technical-debt-related posts on Twitter and find growing interest in the topic, grouped into 11 categories; we cite it as evidence that technical debt is widely discussed in practice, beyond the research literature.

For the SonarQube corroboration subsection, Amanatidis et al. [15] build an empirical benchmark to evaluate how far different technical-debt measurement tools (including SonarQube) agree with each other, showing that the estimate obtained depends on the tool; this contextualises the discrepancy we observe between `jscpd` and SonarQube's duplication detectors as an instance of a documented pattern, not an anomaly specific to our sample or setup. Lenarduzzi et al. [16] use data-driven regression models to estimate the additional project cost of sub-optimal code conditions and report that SonarQube's debt estimates can be made more precise; their work suggests both that automated TD estimation is an active research problem and that the tool's absolute figures are a starting point and not a ground truth, two assumptions that underpin our decision to use SonarQube as a corroboration instrument rather than as the primary measurement.

3 Methodology

3.1 Sample construction

Candidate academic repositories were identified via the GitHub search API using Spanish/Portuguese academic keywords (*tesis*, *trabajo de titulación*, *proyecto integrador*, *trabalho de conclusão*, *tcc*, and related terms) combined with a language filter (Java, Python, JavaScript) and a minimum repository size of 200 KB (GitHub's size figure, which counts all files, not only source code), then verified by checking whether at least one of the repository's top contributors declares a Latin American city or country in their public GitHub profile location field. This combined keyword-plus-location approach is a heuristic, not a ground-truth label, and is discussed as a limitation below. Repositories with fewer than 10 commits were excluded. The GitHub search API excludes forked repositories by default unless a `fork:` qualifier is explicitly added, and none of our queries added one; we verified this by re-issuing representative academic- and control-sample queries and confirming `fork=false` for every returned repository, so the sample does not require a separate fork-filtering step. From 5,418 unique keyword-matched candidates, 300 were accepted as academic (274 successfully cloned and analyzed; the remainder had since been deleted, renamed, or were otherwise unreachable). A control sample of 300 non-academic repositories was drawn with one GitHub search per accepted academic repository: same language, repository size within ±50% of the academic repository's size, results ordered by last update (not by popularity, since an earlier version of this pipeline that sorted by star count pulled in large, well-known projects) and taken from a randomly chosen results page; repositories whose description contained one of the academic keywords were skipped (275 successfully analyzed). The control sample was not restricted to Latin American contributors and was not filtered by number of commits, so the two groups differ in more than their academic label: the comparison is between Latin American academic repositories and size- and language-matched GitHub repositories from any region (Limitations). Each control repository was drawn for one academic repository, but because clone failures broke many of those pairs, the analysis compares the two groups as independent samples, and we refer to the design as a matched-sample comparison, not a paired one.

Table 1 reports the country and language composition of the 300 academic repositories, resolved from the GitHub profile location of each repository's top contributors (the same field used for the LatAm-verification heuristic above). Brazil accounts for nearly half the sample (47.7%), consistent with Brazil's outsized share of Latin American GitHub activity; the remaining sample spans 15 further countries (16 in total).

3.2 Metrics

The primary metrics come from lightweight, license-compatible tools that jointly cover the target metric suite: `lizard` for cross-language cyclomatic complexity [17] and lines of code, `jscpd` for cross-language code duplication, and `ck` (a standard standalone reference tool that runs from a single `.jar` without requiring a SonarQube server) for the Chidamber and Kemerer [7] object-oriented metrics suite on the Java subset (coupling between objects, CBO; lack of cohesion, LCOM; depth of inheritance tree, DIT; weighted methods

Table 1: Composition of the initially accepted academic sample ($n = 300$) by country and language

Country	n	%	Java	JavaScript	Python	Analyzed*
Brazil	143	47.7	57	42	44	130
Argentina	50	16.7	15	25	10	46
Colombia	33	11.0	14	9	10	30
Peru	17	5.7	7	8	2	15
Ecuador	13	4.3	8	2	3	12
Mexico	12	4.0	1	8	3	11
Chile	8	2.7	3	1	4	7
Venezuela	5	1.7	0	4	1	5
El Salvador	4	1.3	1	2	1	4
Uruguay	4	1.3	2	2	0	4
Costa Rica	3	1.0	1	1	1	3
Other	7	2.3	1	4	2	6
Unresolved	1	0.3	1	0	0	1

Other: Bolivia, Cuba, Guatemala, Honduras, Panama; each with fewer than 3 repositories. *26 repositories (8.7%) were unavailable at clone time; the analyzed subset ($n = 274$) retains approximately this proportional distribution. Unresolved: the repository passed the location check at collection time, but its contributors' profile locations no longer named a country when this table was compiled; it is kept in the sample.

per class, WMC). Test coverage was not measured by executing test suites, which is infeasible across 600 arbitrary repositories without individualized build configuration; instead we report the proportion of source files that are test files by naming convention, an explicit proxy rather than measured coverage.

3.3 SonarQube corroboration

As a second instrument we ran SonarQube Community (standalone distribution, Java 21) on a fresh clone of the sample, since the clones used for the primary metrics had been deleted after extraction; 546 of the original 600 repositories could be re-cloned (274 academic, 272 control; a handful had since been deleted, renamed, or made private). SonarQube's Java analyzer refuses to run its full rule set on source files without matching compiled bytecode (`sonar.java.binaries`); since compiling 300-plus arbitrary student and hobby projects is infeasible without per-project build configuration, we pointed this property at an empty directory, which lets analysis proceed with Java-bytecode-dependent rules disabled while complexity, duplication, and code-smell detection remain intact. Under this configuration, 415 of the 546 re-cloned repositories (223 academic, 192 control) completed analysis successfully within a 150-second per-repository budget; the rest failed on very large repositories, unsupported content, or transient errors.

We treat this SonarQube-derived sample as a cross-tool consistency check on the `lizard`/`ck`-based findings above, not a replacement for them (it is smaller, and not repository-for-repository identical to the primary sample), comparing groups with the same Mann-Whitney/Holm/Cliff's-delta procedure on SonarQube-native metrics: cyclomatic and cognitive complexity per function, `duplicated_lines_density`, `sqale_debt_ratio` (SonarQube's own technical-debt-ratio metric, the specific metric our original substitution could not reproduce), and code smells per 1,000 lines.

3.4 Statistical analysis

Because software metrics are typically strongly right-skewed, we compare groups with the Mann-Whitney U test on each of nine metrics, apply the Holm correction for multiple comparisons, and report Cliff's delta [18] as a non-parametric effect size, following standard practice for this kind of comparison in the empirical software engineering literature.

4 Results

To clarify the attrition across analysis stages, Table 2 summarizes the sample size available for each comparison. The primary analysis evaluates the 549 repositories that completed the initial clone and metric-extraction step.

Table 3 summarizes all nine metrics for the primary sample.

Three of the nine metrics remain significant after Holm correction (Figure 1): the two cyclomatic-complexity measures and inheritance depth. All three concern complexity or design structure, but only

Table 2: Sample flow and attrition

Stage	Repositories
Candidates (by keyword)	5,418
Accepted as academic	300
Control sample	300
Successfully cloned and analyzed (Primary)	**549** (274 academic + 275 control)
Re-cloned for SonarQube	546 (274 academic + 272 control)
Completed SonarQube analysis (Sub-sample)	**415** (223 academic + 192 control)

Table 3: Academic vs. control comparison, nine metrics (Holm-corrected). IQR = interquartile range (Q1–Q3)

Metric	Median (acad.)	IQR (acad.)	Median (ctrl.)	IQR (ctrl.)	Cliff's δ	Effect	p_{Holm}
Cyclomatic complexity (mean)	2.16	1.66–3.06	3.12	2.17–4.29	−0.37	medium	< 0.001
Cyclomatic complexity (max)	15.5	9.0–28.0	26.0	15.0–46.0	−0.31	small	< 0.001
DIT (inheritance depth, Java)	1.21	1.06–1.41	1.07	1.00–1.28	+0.29	small	0.003
Lines of code	1,592	592–4,033	2,092	593–6,800	−0.10	negligible	0.322
CBO (coupling, Java)	5.74	3.46–7.12	5.17	1.25–7.56	+0.15	small	0.350
LCOM (cohesion, Java)	8.30	2.61–17.10	7.19	0.46–18.50	+0.12	negligible	0.562
WMC (complexity/class, Java)	5.92	4.34–9.35	6.50	5.04–9.84	−0.08	negligible	0.926
Duplication (%)	5.51	1.96–11.54	5.05	1.26–14.06	+0.03	negligible	1.000
Test-file proportion	0.00	0.00–0.06	0.00	0.00–0.10	−0.01	negligible	1.000

inheritance depth moves in the hypothesized direction. Cyclomatic complexity is *significantly lower* in academic repositories than in the control group, both on average (Cliff's $\delta = -0.37$, a medium effect) and at the maximum observed per repository ($\delta = -0.31$, small), contradicting our initial hypothesis that academic code would be more complex. Inheritance depth (DIT, Java subset only) is significantly higher in academic code ($\delta = +0.29$, small), consistent with the hypothesized pattern of less experienced object-oriented design. Coupling (CBO) trends in the hypothesized direction but does not survive correction ($p = 0.350$). Duplication and test-file proportion show essentially no difference between groups ($\delta \approx 0.03$ and -0.01 respectively), and neither duplication nor size-driven variation in it separates the two groups (Figure 3); notably, the median test-file proportion is *zero in both groups* (Figure 2), meaning near-absent testing is not a distinguishing feature of academic code in this sample; it characterizes both groups about equally.

Because the control group is matched to the academic sample by language as well as by size, and cyclomatic complexity itself varies systematically by programming language, we checked whether the lower-complexity finding is an artifact of language composition. Table 4 repeats the cyclomatic-complexity comparison separately within each of the three languages. The effect holds, and remains significant after a separate Holm correction across the three within-language tests, in all three: Java ($\delta = -0.41$), JavaScript ($\delta = -0.33$), and Python ($\delta = -0.49$), each with medium-to-large effect size. Language composition does not explain the aggregate result. Maximum complexity per repository grows with the number of functions, and the control repositories are somewhat larger (median 180 vs. 153 functions), so we also adjusted that comparison for size: in an OLS regression of log(max CC) on group, the academic coefficient stays negative and significant with log(NLOC) as a covariate ($\beta = -0.33$, $p < 10^{-5}$) and with the log number of functions ($\beta = -0.40$, $p < 10^{-7}$).

Table 4: Cyclomatic complexity (mean), academic vs. control, stratified by language (Holm-corrected across the three within-language tests). The n exclude 10 academic and 12 control repositories in which `lizard` found no measurable function

Language	n (academic/control)	Median (academic)	Median (control)	Cliff's δ	p (Holm)
Java	102 / 96	1.67	2.36	−0.41	< 0.001
JavaScript	93 / 95	2.47	3.08	−0.33	< 0.001
Python	69 / 72	3.20	4.53	−0.49	< 0.001

We also tested a second alternative explanation: that shorter-lived, less-maintained repositories, plausible for course projects that end when a semester does, explain the lower complexity instead of any academic/non-academic distinction. Using each repository's GitHub-reported creation and last-push dates (523 of 549

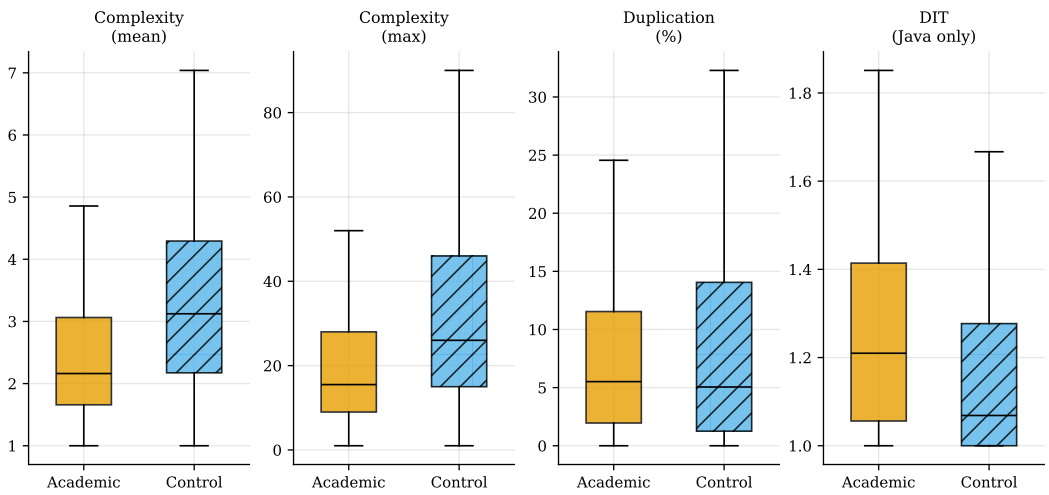


Figure 1: Distribution of key metrics, academic vs. control repositories. Boxes span the interquartile range with the median marked; whiskers extend to 1.5 times the interquartile range; outliers are not shown

repositories resolved; the rest lacked a retrievable date), academic repositories are, if anything, longer-lived than control repositories (median 103.6 days versus 55.2 days, Mann-Whitney $p < 0.001$), the opposite of what the maturity hypothesis predicts. An OLS regression of mean cyclomatic complexity on group membership and log-transformed repository lifespan leaves the group coefficient essentially unchanged from the unadjusted comparison (-0.997 unadjusted versus -0.975 adjusted, both $p < 0.001$) while lifespan itself is not a significant predictor ($p = 0.27$). Repository age does not explain the complexity difference.

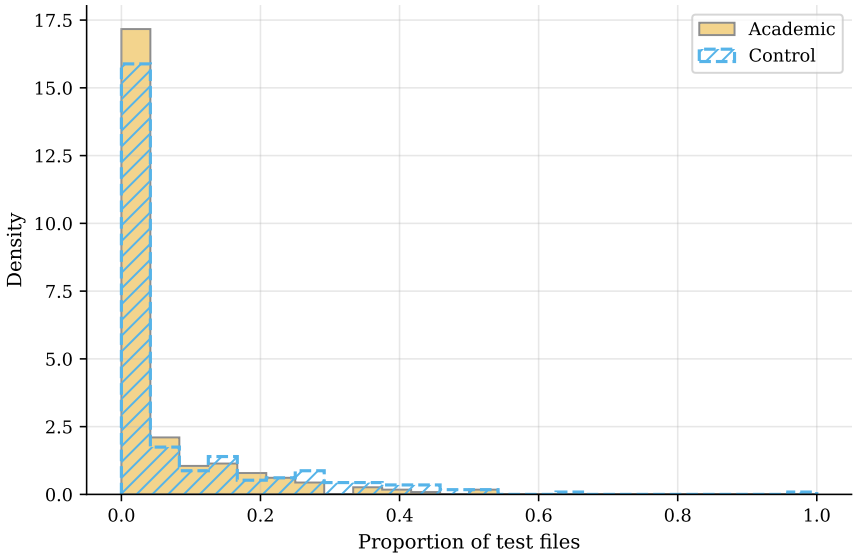


Figure 2: Distribution of the test-file-proportion proxy, academic vs. control

4.1 SonarQube corroboration

SonarQube's own complexity metrics corroborate the `lizard`-based finding across tool architectures: academic repositories have significantly lower complexity per function, both cyclomatic (median 2.07 vs. 2.98, $\delta = -0.458$, medium effect, $p < 10^{-14}$) and cognitive (median 1.31 vs. 3.00, $\delta = -0.426$, medium, $p < 10^{-12}$), the same direction and a comparable effect size to the `lizard`-based result in Table 3. `sqale_debt_ratio`, SonarQube's own technical-debt-ratio metric that our original substitution could not reproduce, is significantly *lower* in academic code ($\delta = -0.142$, $p = 0.025$), consistent with, not contradicting, the lower-complexity result. Duplication tells a more mixed story: SonarQube's duplicate-code detector finds significantly *higher* duplication in academic code (median 2.10% vs. 0.80%, $\delta = +0.181$, small effect, $p = 0.003$), the opposite direction from `jscpd`'s null result in Table 3 ($\delta \approx 0.03$); the Discussion separates the roles of

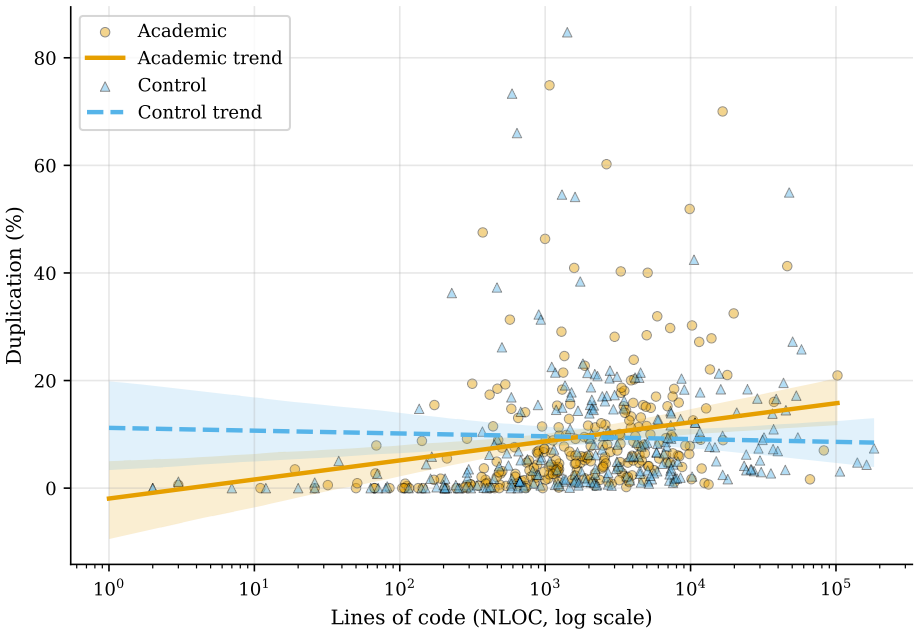


Figure 3: Repository size (lines of code) vs. code duplication, academic vs. control. Trend lines are least-squares fits of duplication on $\log_{10}$(NLOC) for academic (solid) and control (dashed) repositories, with shaded 95% bootstrap confidence bands. Some repositories have fewer than 10 lines in the three analyzed languages (11 academic, 16 control), because the 200 KB size threshold counts all files

Table 5: Academic vs. control comparison on SonarQube-native metrics, 415 repositories (Holm-corrected)

Metric	Median (acad.)	Median (ctrl.)	Cliff's δ	Effect	p_{Holm}
Complexity / function (cyclomatic)	2.07	2.98	−0.458	medium	< 0.001
Complexity / function (cognitive)	1.31	3.00	−0.426	medium	< 0.001
Duplicated lines (%)	2.10	0.80	+0.181	small	0.003
`sqale_debt_ratio` (%)	0.50	0.70	−0.142	negligible	0.025
Code smells / 1,000 lines	27.2	23.3	+0.035	negligible	0.541

sample and detection threshold in this difference. Code smells per 1,000 lines do not differ between groups ($p = 0.541$).

Figure 4 shows the ten most frequent SonarQube code-smell types, normalized per 1,000 lines of code of each group. The most common smell in both groups is a JavaScript/React-specific rule (*prop-types validation*, 3.4 and 3.3 per 1,000 lines), which reflects the language mix of the sample and not an academic-specific pattern. Academic code shows a higher rate of two style-convention smells (`let/const` usage, 1.25 vs. 0.18; package-naming convention, 1.14 vs. 0.08 per 1,000 lines), consistent with less mature tooling and linting setup and not with deeper structural problems. The default-unnamed-package smell (Java) goes the other way (0.11 academic vs. 2.38 control per 1,000 lines). These rates pool each group's lines of code, and we did not store counts by rule and repository, so we cannot tell whether a few repositories dominate a given rule; we therefore do not read the split as uniform across the ten smell types. Finally, we note that the SonarQube sub-sample suffers from differential attrition; see Limitations (§6) for a full accounting of its impact.

4.2 Self-admitted technical debt

As a complement to the structural metrics above, we scanned source comments in the SonarQube sub-sample (223 academic and 191 control repositories with source available; one control repository's clone was missing) for self-admitted technical debt (SATD) markers (`TODO`, `FIXME`, `HACK`, `XXX`, `BUG`, and literal mentions of "technical debt"), normalized per 1,000 lines of code. The median is zero in both groups (most repositories, academic and control alike, contain no SATD marker at all), but the distributions differ: a significantly larger share of academic repositories contain at least one SATD comment (42.2% vs. 23.0%), and academic repositories rank higher overall on this measure (Mann-Whitney $U = 25{,}349.5$, $p < 0.001$, Cliff's $\delta = +0.190$,

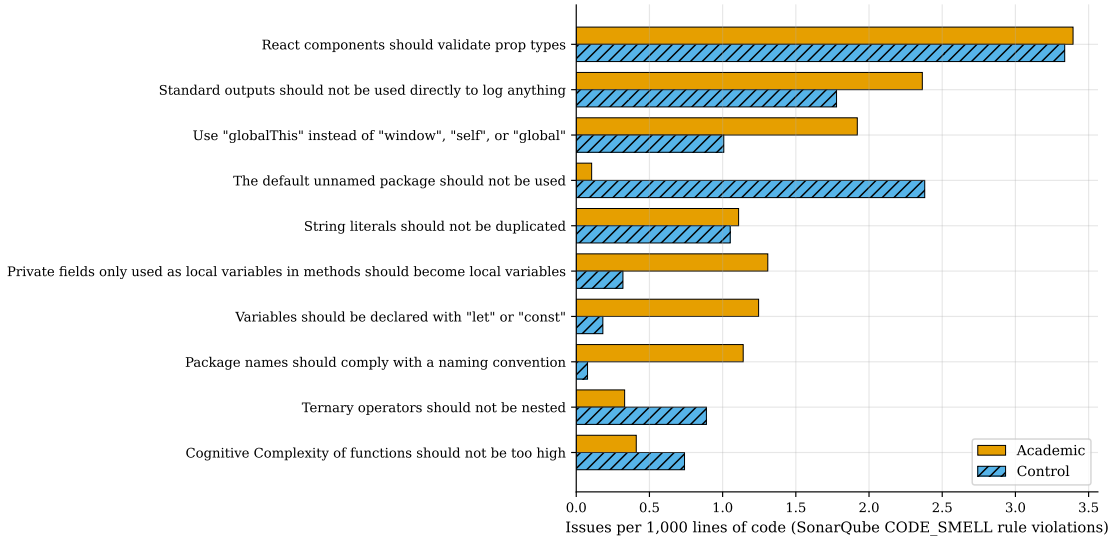


Figure 4: The ten most frequent SonarQube code-smell types, academic vs. control, as issues per 1,000 lines of code of each group (415 repositories analyzed: 223 academic with 1.08 million lines, 192 control with 0.98 million)

small effect; IQR 0.00–0.52 academic vs. 0.00–0.00 control). This difference reflects typical repositories, not extremes: mean SATD density is actually higher in the control group (1.20 vs. 0.65 per 1,000 lines), because a handful of control repositories carry extreme outlier values (up to 120.9 per 1,000 lines) that Cliff's delta, an ordinal measure, is not sensitive to. Unlike the complexity result, this SATD finding runs *with* the original hypothesis: academic developers more often leave explicit, self-admitted markers of incomplete or provisional work in their code than the control group does, even though the code that remains is structurally simpler by the metrics above.

5 Discussion

On this sample, academic Latin American repositories are structurally *simpler*, not more complex, than the size-matched control group. One explanation we considered is a property of the control group, not of the academic group: because repositories were matched to the academic sample's typically modest size, the control sample necessarily excludes large, mature, well-known open-source projects and instead draws from the same broad population of small, low-visibility GitHub repositories that academic work is drawn from: personal projects, bootcamp exercises, and other non-professional code that is not actually representative of production software engineering practice. If the "non-academic" baseline is itself closer to informal, learning-stage code than to professionally maintained software, the comparison this paper runs is better described as *labeled academic vs. unlabeled small repositories of unknown provenance* than as *academic vs. professional*, which would materially change what a lower complexity score in the academic group means: it may reflect course-scoped assignments that are deliberately narrow in requirements, and not academic authors writing better-structured code than typical small non-academic projects. We could test this explanation only crudely. If the control group's low maturity drove the result, control repositories with signs of maturity should be more complex; using stars as that signal, they are less complex instead (median mean CC 2.16 for the 17 control repositories with at least five stars vs. 3.21 for the rest, $p = 0.050$; Spearman $\rho = -0.14$ between stars and complexity, $p = 0.026$). Stars are a weak maturity signal and few control repositories have them, so this does not refute the explanation, but it does not support it either, and the inverted complexity result remains unexplained. A curated control sample of engineered projects would be needed to settle it.

The DIT result, that academic code has measurably deeper inheritance hierarchies, is consistent with a specific, plausible mechanism: object-oriented design courses often assign or scaffold class-hierarchy exercises (e.g., requiring the use of inheritance to satisfy a rubric) in a way that unscaffolded personal projects do not, which would produce deeper hierarchies without those hierarchies reflecting either more sophisticated or more debt-laden design.

We take the near-zero test-file proportion in both groups, not only in the academic group, at face value: it suggests that the near-absence of testing culture in this size class of repository is a broader phenomenon than an academic-specific one, at least by this proxy metric, and that interventions aimed at improving

testing practice in student software may be addressing a gap that is not actually specific to the classroom.

For software engineering teaching in the region, the results point to specific course-level habits, not to a general deficit in student code: inheritance used because an assignment requires it, almost no automated tests, and debt markers left in comments. Active-learning formats used in software engineering courses, among which a recent review in this journal finds games and project-based learning the most frequent, with team-based learning adopted in the authors' own department [19], are natural places to make design review, a minimum of automated tests, and the clean-up of `TODO`/`FIXME` markers explicit deliverables. Because many of these habits carry over into industry, work on training software development teams in the region, which adapts teaching strategies to the learning style most common among practising developers [20], is a complementary route for the same practices after graduation.

That the SonarQube sub-sample reproduces the complexity result with a different parser makes it unlikely to be an artifact of `lizard`'s counting conventions, although both tools measure the same cyclomatic-complexity construct and the sub-sample is affected by differential attrition (Limitations), so this is a consistency check and not an independent replication. The lower `sqale_debt_ratio` in academic code is consistent with this: SonarQube's debt ratio aggregates many rule violations, beyond just complexity, so a lower ratio despite a higher duplication rate suggests that whatever is driving academic code's higher SonarQube-detected duplication is not enough, on its own, to outweigh the complexity advantage in SonarQube's own composite accounting.

The duplication discrepancy between `jscpd` on the primary sample (null) and SonarQube's detector on the 415-repository sub-sample (small, significant, academic higher) mixes two differences: the sample and the detection threshold. We separated them by running `jscpd` on the identical 415-repository sub-sample twice. With its default minimum block of 50 tokens it finds no significant difference (median 4.69% vs. 3.16%, Cliff's $\delta = +0.103$, $p = 0.069$); with a stricter minimum of 100 tokens, which requires longer duplicated blocks, it finds higher duplication in academic code (median 1.77% vs. 1.10%, $\delta = +0.122$, $p = 0.031$), the same direction as SonarQube. Holding the sample fixed, the threshold changes the significance of the result, while the sign of δ is positive at both settings. SonarQube's own copy-paste detector uses a different criterion for Java (at least 10 successive duplicated statements, whatever the number of tokens) and 100 tokens over at least 10 lines for the other languages, so no single `jscpd` setting reproduces it exactly; the remaining gap to SonarQube's effect size ($\delta = +0.181$) is consistent with that difference in how duplicated blocks are defined.

6 Limitations

The keyword-plus-contributor-location heuristic used to identify academic repositories is imperfect and can admit false positives (a repository containing the word "tesis" in an unrelated context) or miss true academic repositories that do not use any of the searched terms; we did not attempt manual verification at scale. The count of 5,418 keyword-matched candidates is the figure reported by our search script at collection time; the unfiltered candidate list was not archived, so that number cannot be re-derived from the released data, whereas the 300 accepted academic and 300 control repositories, and every metric computed on them, can. More generally, GitHub-mining studies are known to over-represent personal, inactive, and non-engineered projects unless explicitly filtered [8, 9]; our size threshold, minimum-commit filter, and (for the academic sample) contributor-location verification each narrow the sample somewhat in this direction, but we did not apply a dedicated engineered-project classifier, so some informal or abandoned repositories likely remain in both groups. The heuristic also identifies only GitHub repositories, excluding GitLab, institutional repositories, and any academic work never pushed to a public host at all, a substantial and likely non-random exclusion, since the most informal, least-polished student work is probably the least likely to be published publicly in the first place. As discussed above, the control group's size-matching means it should not be read as representative of professional software in general. The control group was also not restricted to Latin American contributors and was not filtered by number of commits, and its academic exclusion relied on repository descriptions only; regional origin, minimum activity, and academic status are therefore not held equal between the groups, and part of the differences may reflect them. Keyword-based SATD counts also depend on language and tooling conventions: some IDE templates insert `TODO` comments automatically, which would inflate counts in projects created from them. Test coverage is a naming-convention proxy, not measured execution coverage. `ck`'s object-oriented metrics apply only to the Java subset of each sample (roughly one third), reducing statistical power for CBO, LCOM, DIT, and WMC relative to the language-agnostic metrics. The SonarQube robustness check (Table 5) covers a smaller, non-identical sub-sample (415 of the original 546 re-cloned repositories) because source code had to be re-cloned from GitHub after the original clones were deleted to save disk space, and because SonarQube's Java analyzer, run without compiled bytecode, fails outright on some repositories instead of degrading gracefully. This attrition is not evenly distributed across groups: 223 of 274 academic repositories (81.4%) completed analysis versus 192 of 272 control

repositories (70.6%), an 11-percentage-point gap, and failures skewed toward very large repositories, which are more common in the control group. We checked this by joining the primary sample's `lizard` complexity measurements to SonarQube-attrition status, attrited repositories are far larger than retained ones in both groups (median NLOC 5,497 vs. 1,229 academic, 7,709 vs. 1,369 control, $p < 0.001$ both), which confirms size-driven analyzer failure as the mechanism. On complexity itself, attrited repositories trend higher than retained ones in the academic group (median `cc_mean` 2.35 vs. 2.15) and lower in the control group (2.66 vs. 3.23), a pattern that, if it held up statistically, would mean the SonarQube sub-sample loses relatively complex academic repositories and relatively simple control repositories, which would widen the apparent academic-versus-control complexity gap in Table 5 relative to the true populations, in the same direction as the paper's central finding. However, neither within-group difference reaches significance at these sample sizes ($p = 0.34$ academic, $p = 0.12$ control), so this remains an inconclusive pattern, and we do not treat the SonarQube corroboration as unbiased in its own right. Attrition here is counted against the primary sample (52 academic and 86 control repositories absent from the SonarQube results, the latter including the 3 that could not be re-cloned); the complexity comparison uses the 48 and 78 of them in which `lizard` found measurable functions. `sqale_debt_ratio` and the code-smell-type facets are SonarQube's default rule profile ("Sonar way"), not a profile tuned for either language mix or academic code specifically, and a different quality profile could shift the exact smell-type ranking in Figure 4 without necessarily changing the aggregate complexity and debt-ratio results. Each control repository was drawn for one academic repository, but the pairs were not kept in the analysis (clone failures broke many of them); Mann-Whitney U, an independent-samples test, is therefore the test that matches the analysis we ran, and we refer to the design as matched-sample, not paired. The self-admitted technical debt (SATD) measure is a simple keyword count over source comments, not a validated SATD classifier; it covers only the 415-repository SonarQube sub-sample (not the full 549-repository primary sample, for the same disk-space reason discussed above), and it does not distinguish a genuine debt marker from an unrelated use of a common word (e.g. `BUG` inside an unrelated identifier or string), so the reported counts should be read as a coarse proxy, not a precise measurement. Lines of code is used both as a matching criterion (via the size band) and as one of the nine compared metrics (Table 3); its own comparison ($p = 0.322$, not significant after correction) should be read as a check that the size-band matching worked as intended, not as an independent finding about repository size.

7 Conclusion

In a matched-sample comparison of 274 Latin American academic repositories against 275 size- and language-matched non-academic repositories, academic code shows significantly *lower* cyclomatic complexity than the control group (the opposite of our initial hypothesis) while showing significantly deeper inheritance hierarchies, as hypothesized. Coupling, cohesion, duplication, and a test-coverage proxy show no significant difference after correction, and testing is near-absent in both groups, not distinctively academic. We do not have a confirmed explanation for the inverted complexity result: the composition of the size-matched control group, drawn from small, often informal repositories, is a plausible candidate, but a crude maturity check within the control group does not support it. We recommend that future comparative studies of academic code quality either use a professionally curated control set or explicitly frame their comparison as academic-versus-small-repository, not academic-versus-professional. A check with SonarQube Community on a 415-repository sub-sample reproduces the direction of the lower-complexity finding using SonarQube's own complexity metric and shows academic code has a lower SonarQube technical-debt ratio; because SonarQube and our primary tool compute the same underlying cyclomatic-complexity construct, we treat this as a cross-tool consistency check, not an independent replication, and the one clear point of disagreement across tools, duplication rate, depends on the detection threshold: on the same 415-repository sub-sample, `jscpd` finds no significant difference at its default threshold and higher duplication in academic code at a stricter one, the direction SonarQube reports, with a smaller effect size. On the same sub-sample, a self-admitted technical debt (SATD) comment count runs in the direction the original hypothesis expected, unlike the structural metrics: academic repositories more often contain explicit `TODO`/`FIXME`-style markers of unfinished work (Cliff's $\delta = +0.190$, $p < 0.001$), even though the structural complexity of that code is lower.

Data Availability

The repository lists, extracted metrics, and analysis code are permanently archived and openly accessible in Zenodo at `https://doi.org/10.5281/zenodo.22907773` and GitHub at `https://github.com/JuanAvilesP3/cleiej-deuda-tecnica`. Cloned source code itself is not redistributed; only derived metrics and repository identifiers are shared, consistent with each source repository's own license.

Funding

This research received no external funding. It was conducted within the academic research framework of Escuela Superior Politécnica de Chimborazo (ESPOCH).

Conflict of Interest

The authors declare that they have no conflict of interest.

Use of Generative AI

The authors used generative AI tools (large language models) as an assistant for language editing, support in writing and debugging analysis code, and LaTeX formatting. The research questions, study design, data collection, execution of the analyses, results, and their interpretation are the authors' own work; all AI-assisted output was reviewed and verified by the authors, who take full responsibility for the content of this article. No figure was produced with image-generation tools.

Author Contributions

Isaac David Torres-Paredes: Conceptualization, Methodology, Project Administration, Supervision, Writing - Original Draft. **Juan Pablo Aviles-Esparza:** Software, Formal Analysis, Data Curation, Visualization, Writing - Review & Editing. **Italo Javier Tenempaguay-Granizo:** Validation, Investigation, Writing - Review & Editing.

References

[1] W. Cunningham, "The WyCash portfolio management system," OOPSLA'92 Experience Report, Addendum to the Proceedings, Tech. Rep., 1992.

[2] E. Tom, A. Aurum, and R. Vidgen, "An exploration of technical debt," *Journal of Systems and Software*, vol. 86, no. 6, pp. 1498–1516, 2013. [Online]. Available: https://doi.org/10.1016/j.jss.2012.12.052

[3] Z. Li, P. Avgeriou, and P. Liang, "A systematic mapping study on technical debt and its management," *Journal of Systems and Software*, vol. 101, pp. 193–220, 2015. [Online]. Available: https://doi.org/10.1016/j.jss.2014.12.027

[4] D. Santos, A. Resende, P. Afonso Junior, and H. Costa, "External quality metrics for object-oriented software: A systematic literature review," *CLEI Electronic Journal*, vol. 20, no. 3, 2017. [Online]. Available: https://doi.org/10.19153/cleiej.20.3.3

[5] S. Hamer, C. Quesada-López, A. Martínez, and M. Jenkins, "Using git metrics to measure students' and teams' code contributions in software development projects," *CLEI Electronic Journal*, vol. 24, no. 2, 2021. [Online]. Available: https://doi.org/10.19153/cleiej.24.2.8

[6] C. Mateos, A. Zunino, S. Misra, D. Anabalon, and A. Flores, "Managing web service interface complexity via an oo metric-based early approach," *CLEI Electronic Journal*, vol. 20, no. 3, 2017. [Online]. Available: https://doi.org/10.19153/cleiej.20.3.2

[7] S. R. Chidamber and C. F. Kemerer, "A metrics suite for object oriented design," *IEEE Transactions on Software Engineering*, vol. 20, no. 6, pp. 476–493, 1994. [Online]. Available: https://doi.org/10.1109/32.295895

[8] E. Kalliamvakou, G. Gousios, K. Blincoe, L. Singer, D. M. German, and D. Damian, "The promises and perils of mining GitHub," in *Proceedings of the 11th Working Conference on Mining Software Repositories (MSR)*, 2014, pp. 92–101. [Online]. Available: https://doi.org/10.1145/2597073.2597074

[9] N. Munaiah, S. Kroh, C. Cabrey, and M. Nagappan, "Curating GitHub for engineered software projects," *Empirical Software Engineering*, vol. 22, no. 6, pp. 3219–3253, 2017. [Online]. Available: https://doi.org/10.1007/s10664-017-9512-6

[10] A. Potdar and E. Shihab, "An exploratory study on self-admitted technical debt," in *2014 IEEE International Conference on Software Maintenance and Evolution (ICSME)*, 2014, pp. 91–100, [Online], Available: https://doi.org/10.1109/icsme.2014.31.

[11] G. Bavota and B. Russo, "A large-scale empirical study on self-admitted technical debt," in *Proceedings of the 13th International Conference on Mining Software Repositories (MSR)*, 2016, pp. 315–326. [Online]. Available: https://doi.org/10.1145/2901739.2901742

[12] E. d. S. Maldonado and E. Shihab, "Detecting and quantifying different types of self-admitted technical debt," in *Proceedings of the 7th IEEE International Workshop on Managing Technical Debt (MTD)*, 2015, pp. 9–15. [Online]. Available: https://doi.org/10.1109/MTD.2015.7332619

[13] Q. Huang, E. Shihab, X. Xia, D. Lo, and S. Li, "Identifying self-admitted technical debt in open source projects using text mining," *Empirical Software Engineering*, vol. 23, no. 1, pp. 418–451, 2018. [Online]. Available: https://doi.org/10.1007/s10664-017-9522-4

[14] R. Alfayez, R. Winn, Y. Ding, G. Alfayez, and B. Boehm, "Technical debt (TD) through the lens of Twitter: A survey," *Journal of Software: Evolution and Process*, vol. 36, no. 4, p. e2547, 2024, [Online], Available: https://doi.org/10.1002/smr.2547.

[15] T. Amanatidis, N. Mittas, A. Moschou, A. Chatzigeorgiou, A. Ampatzoglou, and L. Angelis, "Evaluating the agreement among technical debt measurement tools: Building an empirical benchmark of technical debt liabilities," *Empirical Software Engineering*, vol. 25, no. 5, pp. 4161–4204, 2020. [Online]. Available: https://doi.org/10.1007/s10664-020-09869-w

[16] V. Lenarduzzi, A. Martini, D. Taibi, and D. A. Tamburri, "Towards surgically-precise technical debt estimation: Early results and research roadmap," in *Proceedings of the 3rd ACM SIGSOFT International Workshop on Machine Learning Techniques for Software Quality Evaluation (MaLTeSQuE)*, 2019, pp. 37–42. [Online]. Available: https://doi.org/10.1145/3340482.3342747

[17] T. J. McCabe, "A complexity measure," *IEEE Transactions on Software Engineering*, vol. SE-2, no. 4, pp. 308–320, 1976. [Online]. Available: https://doi.org/10.1109/TSE.1976.233837

[18] N. Cliff, "Dominance statistics: Ordinal analyses to answer ordinal questions," *Psychological Bulletin*, vol. 114, no. 3, pp. 494–509, 1993. [Online]. Available: https://doi.org/10.1037/0033-2909.114.3.494

[19] J. I. Irabedra, M. Solari, and G. Mousqués, "Active learning pedagogical strategies in software engineering university courses," *CLEI Electronic Journal*, vol. 28, no. 5, pp. 6:1–6:14, 2025. [Online]. Available: https://doi.org/10.19153/cleiej.28.5.6

[20] M. Hidalgo, K. Rodriguez, L. M. Castro, F. Montoya, and H. Astudillo, "Crosscutting teaching strategies for software development teams: Insights from Kolb's inventory," *CLEI Electronic Journal*, vol. 28, no. 5, pp. 5:1–5:20, 2025. [Online]. Available: https://doi.org/10.19153/cleiej.28.5.5